

■ SQL

```

SELECT C.pid
FROM   Catalog C, Suppliers S
WHERE  S.sname = 'Yosemite Sham' AND C.sid = S.sid
      AND C.cost ≥ ALL (Select C2.cost
                        FROM   Catalog C2, Suppliers S2
                        WHERE  S2.sname = 'Yosemite Sham'
                        AND C2.sid = S2.sid)

```

Exercise 4.4 Consider the Supplier-Parts-Catalog schema from the previous question. State what the following queries compute:

1. $\pi_{sname}(\pi_{sid}(\sigma_{color='red'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers)$
2. $\pi_{sname}(\pi_{sid}((\sigma_{color='red'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers))$
3. $(\pi_{sname}((\sigma_{color='red'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers)) \cap$
 $(\pi_{sname}((\sigma_{color='green'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers))$
4. $(\pi_{sid}((\sigma_{color='red'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers)) \cap$
 $(\pi_{sid}((\sigma_{color='green'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers))$
5. $\pi_{sname}((\pi_{sid,sname}((\sigma_{color='red'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers)) \cap$
 $(\pi_{sid,sname}((\sigma_{color='green'} Parts) \bowtie (\sigma_{cost < 100} Catalog) \bowtie Suppliers))))$

Answer 4.4 The statements can be interpreted as:

1. Find the Supplier names of the suppliers who supply a red part that costs less than 100 dollars.
2. This Relational Algebra statement does not return anything because of the sequence of projection operators. Once the sid is projected, it is the only field in the set. Therefore, projecting on sname will not return anything.
3. Find the Supplier names of the suppliers who supply a red part that costs less than 100 dollars and a green part that costs less than 100 dollars.
4. Find the Supplier ids of the suppliers who supply a red part that costs less than 100 dollars and a green part that costs less than 100 dollars.
5. Find the Supplier names of the suppliers who supply a red part that costs less than 100 dollars and a green part that costs less than 100 dollars.

Exercise 4.5 Consider the following relations containing airline flight information:

```

Flights(flno: integer, from: string, to: string,
        distance: integer, departs: time, arrives: time)
Aircraft(aid: integer, aname: string, cruisingrange: integer)
Certified(eid: integer, aid: integer)
Employees(eid: integer, ename: string, salary: integer)

```

Note that the Employees relation describes pilots and other kinds of employees as well; every pilot is certified for some aircraft (otherwise, he or she would not qualify as a pilot), and only pilots are certified to fly.

Write the following queries in relational algebra, tuple relational calculus, and domain relational calculus. Note that some of these queries may not be expressible in relational algebra (and, therefore, also not expressible in tuple and domain relational calculus)! For such queries, informally explain why they cannot be expressed. (See the exercises at the end of Chapter 5 for additional queries over the airline schema.)

1. Find the *eids* of pilots certified for some Boeing aircraft.
2. Find the *names* of pilots certified for some Boeing aircraft.
3. Find the *aids* of all aircraft that can be used on non-stop flights from Bonn to Madras.
4. Identify the flights that can be piloted by every pilot whose salary is more than \$100,000.
5. Find the names of pilots who can operate planes with a range greater than 3,000 miles but are not certified on any Boeing aircraft.
6. Find the *eids* of employees who make the highest salary.
7. Find the *eids* of employees who make the second highest salary.
8. Find the *eids* of employees who are certified for the largest number of aircraft.
9. Find the *eids* of employees who are certified for exactly three aircraft.
10. Find the total amount paid to employees as salaries.
11. Is there a sequence of flights from Madison to Timbuktu? Each flight in the sequence is required to depart from the city that is the destination of the previous flight; the first flight must leave Madison, the last flight must reach Timbuktu, and there is no restriction on the number of intermediate flights. Your query must determine whether a sequence of flights from Madison to Timbuktu exists for *any* input Flights relation instance.

Answer 4.5 In the answers below RA refers to Relational Algebra, TRC refers to Tuple Relational Calculus and DRC refers to Domain Relational Calculus.

1. ■ RA

$$\pi_{eid}(\sigma_{aname='Boeing'}(Aircraft \bowtie Certified))$$

■ TRC

$$\{C.eid \mid C \in Certified \wedge \\ \exists A \in Aircraft(A.aid = C.aid \wedge A.aname = 'Boeing')\}$$

■ DRC

$$\{\langle C.eid \rangle \mid \langle C.eid, C.aid \rangle \in Certified \wedge \\ \exists Aid, AN, AR(\langle Aid, AN, AR \rangle \in Aircraft \\ \wedge Aid = C.aid \wedge AN = 'Boeing')\}$$

■ SQL

```
SELECT C.eid
FROM   Aircraft A, Certified C
WHERE  A.aid = C.aid AND A.aname = 'Boeing'
```

2. ■ RA

$$\pi_{ename}(\sigma_{aname='Boeing'}(Aircraft \bowtie Certified \bowtie Employees))$$

■ TRC

$$\{E.ename \mid E \in Employees \wedge \exists C \in Certified \\ (\exists A \in Aircraft(A.aid = C.aid \wedge A.aname = 'Boeing' \wedge E.eid = C.eid))\}$$

■ DRC

$$\{\langle EN \rangle \mid \langle Eid, EN, ES \rangle \in Employees \wedge \\ \exists C.eid, C.aid(\langle C.eid, C.aid \rangle \in Certified \wedge \\ \exists Aid, AN, AR(\langle Aid, AN, AR \rangle \in Aircraft \wedge \\ Aid = C.aid \wedge AN = 'Boeing' \wedge Eid = C.eid))\}$$

■ SQL

```
SELECT E.ename
FROM   Aircraft A, Certified C, Employees E
WHERE  A.aid = C.aid AND A.aname = 'Boeing' AND E.eid = C.eid
```

3. ■ RA

$$\rho(\text{BonnToMadrid}, \sigma_{\text{from}='Bonn' \wedge \text{to}='Madrid'}(\text{Flights}))$$

$$\pi_{\text{aid}}(\sigma_{\text{cruisingrange} > \text{distance}}(\text{Aircraft} \times \text{BonnToMadrid}))$$

■ TRC

$$\{A.\text{aid} \mid A \in \text{Aircraft} \wedge \exists F \in \text{Flights}$$

$$(F.\text{from} = 'Bonn' \wedge F.\text{to} = 'Madrid' \wedge A.\text{cruisingrange} > F.\text{distance})\}$$

■ DRC

$$\{Aid \mid \langle Aid, AN, AR \rangle \in \text{Aircraft} \wedge$$

$$(\exists FN, FF, FT, FDi, FDe, FA (\langle FN, FF, FT, FDi, FDe, FA \rangle \in \text{Flights} \wedge$$

$$FF = 'Bonn' \wedge FT = 'Madrid' \wedge FDi < AR))\}$$

■ SQL

```
SELECT A.aid
FROM   Aircraft A, Flights F
WHERE  F.from = 'Bonn' AND F.to = 'Madrid' AND
      A.cruisingrange > F.distance
```

4. ■ RA

$$\pi_{\text{flno}}(\sigma_{\text{distance} < \text{cruisingrange} \wedge \text{salary} > 100,000}(\text{Flights} \bowtie \text{Aircraft} \bowtie$$

$$\text{Certified} \bowtie \text{Employees})))$$

■

$$\text{TRC } \{F.\text{flno} \mid F \in \text{Flights} \wedge \exists A \in \text{Aircraft} \exists C \in \text{Certified}$$

$$\exists E \in \text{Employees} (A.\text{cruisingrange} > F.\text{distance} \wedge E.\text{salary} > 100,000 \wedge$$

$$A.\text{aid} = C.\text{aid} \wedge E.\text{eid} = C.\text{eid})\}$$

■ DRC

$$\{FN \mid \langle FN, FF, FT, FDi, FDe, FA \rangle \in \text{Flights} \wedge$$

$$\exists C.\text{eid}, C.\text{aid} (\langle C.\text{eid}, C.\text{aid} \rangle \in \text{Certified} \wedge$$

$$\exists A.\text{aid}, AN, AR (\langle A.\text{aid}, AN, AR \rangle \in \text{Aircraft} \wedge$$

$$\exists E.\text{eid}, EN, ES (\langle E.\text{eid}, EN, ES \rangle \in \text{Employees}$$

$$(AR > FDi \wedge ES > 100,000 \wedge A.\text{aid} = C.\text{aid} \wedge E.\text{eid} = C.\text{eid}))\}$$

■ SQL

```
SELECT E.ename
FROM   Aircraft A, Certified C, Employees E, Flights F
WHERE  A.aid = C.aid AND E.eid = C.eid AND
      distance < cruisingrange AND salary > 100,000
```

5. ■ RA $\rho(R1, \pi_{eid}(\sigma_{cruisingrange > 3000}(Aircraft \bowtie Certified)))$
 $\pi_{ename}(Employees \bowtie (R1 - \pi_{eid}(\sigma_{aname='Boeing'}(Aircraft \bowtie Certified))))$

■ TRC

$\{E.ename \mid E \in Employees \wedge \exists C \in Certified(\exists A \in Aircraft$
 $(A.aid = C.aid \wedge E.eid = C.eid \wedge A.cruisingrange > 3000)) \wedge$
 $\neg(\exists C2 \in Certified(\exists A2 \in Aircraft(A2.aname = 'Boeing' \wedge C2.aid =$
 $A2.aid \wedge C2.eid = E.eid)))\}$

■ DRC

$\{\langle EN \rangle \mid \langle Eid, EN, ES \rangle \in Employees \wedge$
 $\exists Ceid, Caid(\langle Ceid, Caid \rangle \in Certified \wedge$
 $\exists Aid, AN, AR(\langle Aid, AN, AR \rangle \in Aircraft \wedge$
 $Aid = Caid \wedge Eid = Ceid \wedge AR > 3000)) \wedge$
 $\neg(\exists Aid2, AN2, AR2(\langle Aid2, AN2, AR2 \rangle \in Aircraft \wedge$
 $\exists Ceid2, Caid2(\langle Ceid2, Caid2 \rangle \in Certified$
 $\wedge Aid2 = Caid2 \wedge Eid = Ceid2 \wedge AN2 = 'Boeing')))\}$

■ SQL

```
SELECT E.ename
FROM   Certified C, Employees E, Aircraft A
WHERE  A.aid = C.aid AND E.eid = C.eid AND A.cruisingrange > 3000
AND E.eid NOT IN ( SELECT C2.eid
FROM   Certified C2, Aircraft A2
WHERE  C2.aid = A2.aid AND A2.aname = 'Boeing' )
```

6. ■ RA

The approach to take is first find all the employees who do not have the highest salary. Subtract these from the original list of employees and what is left is the highest paid employees.

$\rho(E1, Employees)$
 $\rho(E2, Employees)$
 $\rho(E3, \pi_{E2.eid}(E1 \bowtie_{E1.salary > E2.salary} E2))$
 $(\pi_{eid} E1) - E3$

■ TRC

$\{E1.eid \mid E1 \in Employees \wedge \neg(\exists E2 \in Employees(E2.salary > E1.salary))\}$

■ DRC

8. This cannot be expressed in relational algebra (or calculus) because there is no operator to count, and this query requires the ability to count up to a number that depends on the data. The query can however be expressed in SQL as follows:

```

SELECT Temp.eid
FROM   ( SELECT   C.eid AS eid, COUNT (C.aid) AS cnt,
              FROM     Certified C
              GROUP BY C.eid) AS Temp
WHERE  Temp.cnt = ( SELECT   MAX (Temp.cnt)
                   FROM     Temp)

```

9. ■ RA

The approach behind this query is to first find the employees who are certified for at least three aircraft (they appear at least three times in the Certified relation). Then find the employees who are certified for at least four aircraft. Subtract the second from the first and what is left is the employees who are certified for exactly three aircraft.

$$\begin{aligned}
& \rho(R1, \text{Certified}) \\
& \rho(R2, \text{Certified}) \\
& \rho(R3, \text{Certified}) \\
& \rho(R4, \text{Certified}) \\
& \rho(R5, \pi_{\text{eid}}(\sigma_{(R1.\text{eid}=R2.\text{eid}=R3.\text{eid}) \wedge (R1.\text{aid} \neq R2.\text{aid} \neq R3.\text{aid})}(R1 \times R2 \times R3))) \\
& \rho(R6, \pi_{\text{eid}}(\sigma_{(R1.\text{eid}=R2.\text{eid}=R3.\text{eid}=R4.\text{eid}) \wedge (R1.\text{aid} \neq R2.\text{aid} \neq R3.\text{aid} \neq R4.\text{aid})}(R1 \times R2 \times R3 \times R4))) \\
& R5 - R6
\end{aligned}$$

■ TRC

$$\begin{aligned}
& \{C1.\text{eid} \mid C1 \in \text{Certified} \wedge \exists C2 \in \text{Certified} (\exists C3 \in \text{Certified} \\
& (C1.\text{eid} = C2.\text{eid} \wedge C2.\text{eid} = C3.\text{eid} \wedge \\
& C1.\text{aid} \neq C2.\text{aid} \wedge C2.\text{aid} \neq C3.\text{aid} \wedge C3.\text{aid} \neq C1.\text{aid} \wedge \\
& \neg(\exists C4 \in \text{Certified} \\
& (C3.\text{eid} = C4.\text{eid} \wedge C1.\text{aid} \neq C4.\text{aid} \wedge \\
& C2.\text{aid} \neq C4.\text{aid} \wedge C3.\text{aid} \neq C4.\text{aid})))\}
\end{aligned}$$

■ DRC

$$\begin{aligned}
& \{\langle CE1 \rangle \mid \langle CE1, CA1 \rangle \in \text{Certified} \wedge \\
& \exists CE2, CA2 (\langle CE2, CA2 \rangle \in \text{Certified} \wedge \\
& \exists CE3, CA3 (\langle CE3, CA3 \rangle \in \text{Certified} \wedge \\
& (CE1 = CE2 \wedge CE2 = CE3 \wedge \\
& CA1 \neq CA2 \wedge CA2 \neq CA3 \wedge CA3 \neq CA1 \wedge \\
& \neg(\exists CE4, CA4 (\langle CE4, CA4 \rangle \in \text{Certified} \wedge
\end{aligned}$$

$$(CE3 = CE4 \wedge CA1 \neq CA4 \wedge \\ CA2 \neq CA4 \wedge CA3 \neq CA4))))\}$$

■ SQL

```
SELECT C1.eid
FROM   Certified C1, Certified C2, Certified C3
WHERE  (C1.eid = C2.eid AND C2.eid = C3.eid AND
        C1.aid ≠ C2.aid AND C2.aid ≠ C3.aid AND C3.aid ≠ C1.aid)
EXCEPT
SELECT C4.eid
FROM   Certified C4, Certified C5, Certified C6, Certified C7,
WHERE  (C4.eid = C5.eid AND C5.eid = C6.eid AND C6.eid = C7.eid AND
        C4.aid ≠ C5.aid AND C4.aid ≠ C6.aid AND C4.aid ≠ C7.aid AND
        C5.aid ≠ C6.aid AND C5.aid ≠ C7.aid AND C6.aid ≠ C7.aid )
```

This could also be done in SQL using COUNT.

10. This cannot be expressed in relational algebra (or calculus) because there is no operator to sum values. The query can however be expressed in SQL as follows:

```
SELECT SUM (E.salaries)
FROM   Employees E
```

11. This cannot be expressed in relational algebra or relational calculus or SQL. The problem is that there is no restriction on the number of intermediate flights. All of the query methods could find if there was a flight directly from Madison to Timbuktu and if there was a sequence of two flights that started in Madison and ended in Timbuktu. They could even find a sequence of n flights that started in Madison and ended in Timbuktu as long as there is a static (i.e., data-independent) upper bound on the number of intermediate flights. (For large n , this would of course be long and impractical, but at least possible.) In this query, however, the upper bound is not static but dynamic (based upon the set of tuples in the Flights relation).

In summary, if we had a static upper bound (say k), we could write an algebra or SQL query that repeatedly computes (upto k) joins on the Flights relation. If the upper bound is dynamic, then we cannot write such a query because k is not known when writing the query.

Exercise 4.6 What is *relational completeness*? If a query language is relationally complete, can you write any desired query in that language?

Answer 4.6 Answer omitted.